



Case 2 - Execution Plan

Arjun / Kargo

MESA AI AND ITS APPLICATIONS

COHORT C4

The Problem Isn't Too Many CVs

60

Applications Received

Across PM and SPM roles, accumulated
over 11 weeks

19

Opened

The remaining 41 have never been
looked at

0 Offers

No Hiring Decisions Made

Despite consistent application volume,
no hiring decisions made

Kargo cannot hit its December headcount target on this trajectory.

✓ **Goal** → A ranked shortlist per role, with enough reasoning per candidate that Arjun can make a call decision in under 10 minutes.

The Nine Checks - Does This Deserve to Be Built?

Kill Switches

Any No means: don't build it.

- **01 Problem real? YES** - 60 received, 19 opened, 0 offers in 11 weeks. Cost is measurable.
- **02 Workflow repeated? YES** - Applications arrive weekly. The review-and-stall cycle repeats every batch.
- **03 Input available? YES** - 60 CVs, 8 hire profiles, 2 role specs. All inputs exist now.

Sizing

Any No means: build something smaller.

- **04 Output valuable? YES** - Ranked shortlist with scores and briefs replaces 60 unreviewed PDFs.
- **05 Impact measurable? YES** - Baseline: 0 offers in 11 weeks. Target: first offer within 2 weeks.
- **08 ROI worth it? YES** - A PM hire = 6-12 months of product velocity. Every week costs Kargo.

Boundary

Any No means: put a human exactly here.

- **06 Failure risk OK? YES** for ranking and scoring. **NO** for auto-rejection - A wrongly rejected candidate never returns. No feedback loop, no recovery.
- **07 Judgment protected? YES** - Arjun reviews every shortlist and confirms every send.
- **09 Owner clear? YES** for the shortlist. **NO** for auto-rejection - No named person is accountable for a system rejection. That is a liability.

THE CUT

What Arjun Asked For

Auto-send emails to rejected and selected candidates, only after he confirms.

Killed By - Check 06 + Check 09

- Check 06: A wrongly rejected candidate never knows and never returns. The failure is invisible - no feedback loop, no recovery possible.
- Check 09: When the system rejects a candidate, no person can be named as accountable for that specific decision. That is not a workflow. It is a liability.

What Gets Built Instead

What gets automated:


- Extraction, scoring against the rubric, ranking, brief generation, email drafting.

What stays human:

- The shortlist review, the confirm-to-send decision, the hiring call.

The system ranks and explains. You look at everything below the line once. That is ten minutes, not ten hours - because the ranking does the heavy lifting. We do not reject on your behalf because a wrong rejection is one you cannot fix.

Case 2 · Hiring Dashboard — Components Map



The Stack



Claude Chat

Builds the rubric - no code, before anything else is touched.



Supabase

Stores candidate records, scores, and the rubric.



Resend

Delivers emails to candidates at the end of the build.



Claude Code

Where the project lives and gets built end-to-end.



Vercel

Puts the dashboard live online.



Gemini Flash

All four AI steps: extraction, scoring, brief generation, email draft.



GitHub

Saves progress between every session.

Checkpoint L4.1 - Design the Rubric in Claude Chat

1 Open Claude Chat

In the browser - no code required.

2 Upload the hires/ and roles/ folders

8 past hire profiles + 2 role specs (PM and SPM).

3 Paste the outcomes table and send the prompt

Tell Claude the performance rating for each hire. Ask it to find 2–3 patterns in the high-rated hires that do not appear in the job descriptions. Then ask it to design a 4–6 criterion rubric with weights that add to 100% for PM and separately for SPM.

4 Save the result

rubric.txt

✓ **DONE WHEN:** The rubric has 4–6 criteria for each role, weights add to 100%, and you can point to a hire profile behind each criterion - not a job description.

Checkpoint L4.2 - Build the Project in Claude Code

1 Upload rubric.txt into Claude Code

Before typing anything.

2 Send the opening prompt

The key line: "Build the complete project - database, all pages, all scoring and brief and email draft logic, email sending via Resend."

3 Add keys to .env.local

SUPABASE_URL, SUPABASE_ANON_KEY, GEMINI_API_KEY.
Leave RESEND_API_KEY blank for now.

4 Deploy

Push to GitHub → import into Vercel → add environment variables → deploy.

❌ **Security:** .env.local stays private. Confirm it is listed in .gitignore before pushing.

✅ **DONE WHEN:** Vercel shows the project is live. The Supabase database is connected and has the rubric data, from rubric.txt.

The Data Privacy Block

You've built a system that processes people who never applied to Kargo. Let's understand the implication before continuing.

1 What is the difference between the Gemini AI Studio free tier and the Gemini API with billing enabled?

The free tier may use your inputs to improve Google's models. The billed API does not. One configuration line. Entirely different data treatment.

2 What does the extraction step do that makes the pipeline compliant in a real deployment?

It separates personal details from CV content at ingestion. Every AI call after that receives only content — no name, email, or phone. The identifier stays in Supabase and never enters Gemini.

Checkpoint B.1 - Upload CVs and Test the Pipeline

1 Upload 3 test CVs first

Start by uploading one clearly strong PM, one clearly weak SPM, and one ambiguous candidate CV. Remember to select the correct role (PM or SPM) for each upload.

2 Check the extraction in Supabase

Verify that the personal details field (name, email, phone) is visible, but the CV content field contains no sensitive personal details like name, email, or phone. This ensures privacy compliance.

3 Check the scores

Confirm that both PM and SPM rubric scores are generated and exist for every candidate uploaded, reflecting the criteria defined in the rubric.

4 Check the dashboard

Ensure candidates are ranked by their scores on the dashboard. Verify that the top candidates have a concise 3-sentence brief and every candidate has a draft email generated.

5 Upload all 60 CVs

Once the test CVs confirm the pipeline is working as expected, proceed to upload all remaining 60 CVs to process them through the system.

✔ **DONE WHEN:** All 60 rows are populated in Supabase with scores. The top 5 candidates per role have a brief and a draft invite. Every other candidate has a draft rejection email.



Checkpoints B·2 + B·3 - Resend and End-to-End Test

B·2 Resend Integration

- Create a free account at resend.com
- Generate your API key from the Resend dashboard
- Add `RESEND_API_KEY` to Vercel environment variables
- Redeploy
- Tell Claude Code to wire the Confirm button to Resend – sending to the MESA test address for this session

✅ **DONE WHEN:** Clicking Confirm triggers Resend. Test inbox receives the email within 30 seconds. Candidate's real name shows, not `[NAME]`.

B·3 Review UI and End-to-End Test

- Upload 2 new CVs
- Read the top candidate's brief
- Read the draft email
- Click Confirm
- Check the MESA test inbox
- Confirm sent status updates in Supabase

✅ **DONE WHEN:** The full loop - upload to send - runs in under 5 minutes. Push to GitHub.